

CENG443

Heterogeneous Parallel Programming

Parallelism

The summary of the concepts in the book:

Computer Architecture, A Quantitative Approach. John L. Hennessy, David A. Patterson

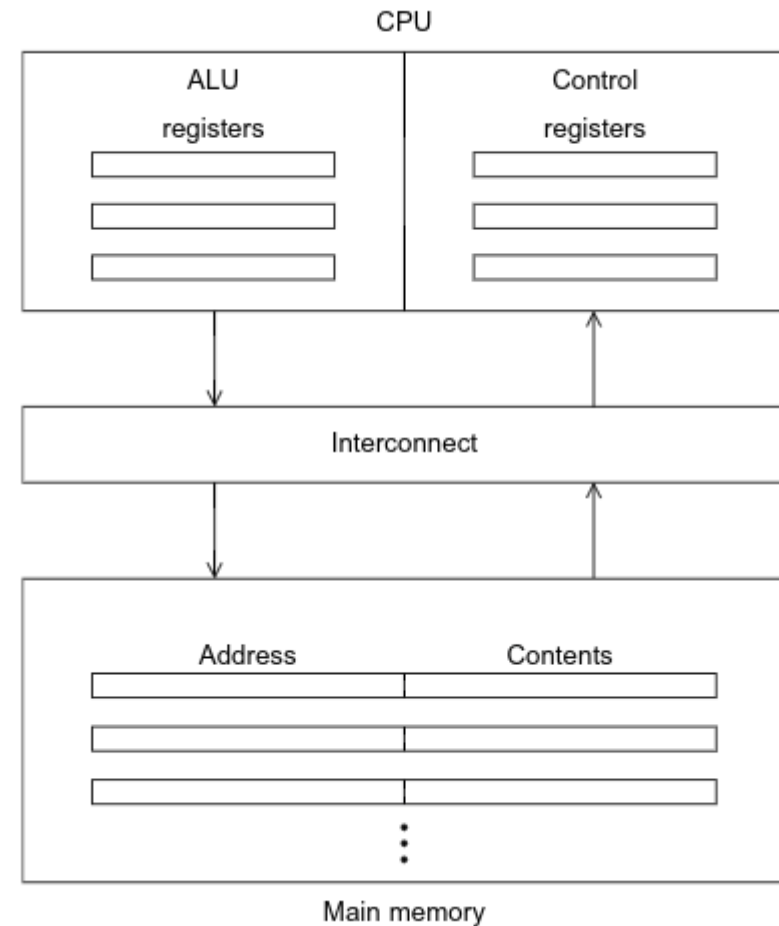


The von Neumann Architecture

Main memory

**Central processing unit
(CPU)**


Interconnection



Main Memory

Collection of locations, each of which is capable of storing both instructions and data

Every location consists of an address, which is used to access the location, and the contents of the location




Central Processing Unit (CPU)

Control unit : responsible for deciding which instruction in a program should be executed

Arithmetic and logic unit (ALU) : responsible for executing the actual instructions

Register : quickly accessible location available to CPU

memory




fetch/read



CPU

memory



write/store



CPU

Example Program

Compute $\sin(x)$ using taylor expansion: $\sin(x) = x - x^3/3! + x^5/5! + x^7/7! + \dots$

```
void sinx(int N, int terms, float* x, float* result)
{
    for (int i=0; i<N; i++)
    {
        float value = x[i];
        float numer = x[i] * x[i] * x[i];
        int denom = 6; // 3!
        int sign = -1;

        for (int j=1; j<=terms; j++)
        {
            value += sign * numer / denom
            numer *= x[i] * x[i];
            denom *= (j+3) * (j+4);
            sign *= -1;
        }

        result[i] = value;
    }
}
```

Compile Program

```
void sinx(int N, int terms, float* x, float* result)
{
    for (int i=0; i<N; i++)
    {
        float value = x[i];
        float number = x[i] * x[i] * x[i];
        int denom = 6; // 3!
        int sign = -1;

        for (int j=1; j<=terms; j++)
        {
            value += sign * number / denom;
            number *= x[i] * x[i];
            denom *= (j+3) * (j+4);
            sign *= -1;
        }

        result[i] = value;
    }
}
```

x[i]

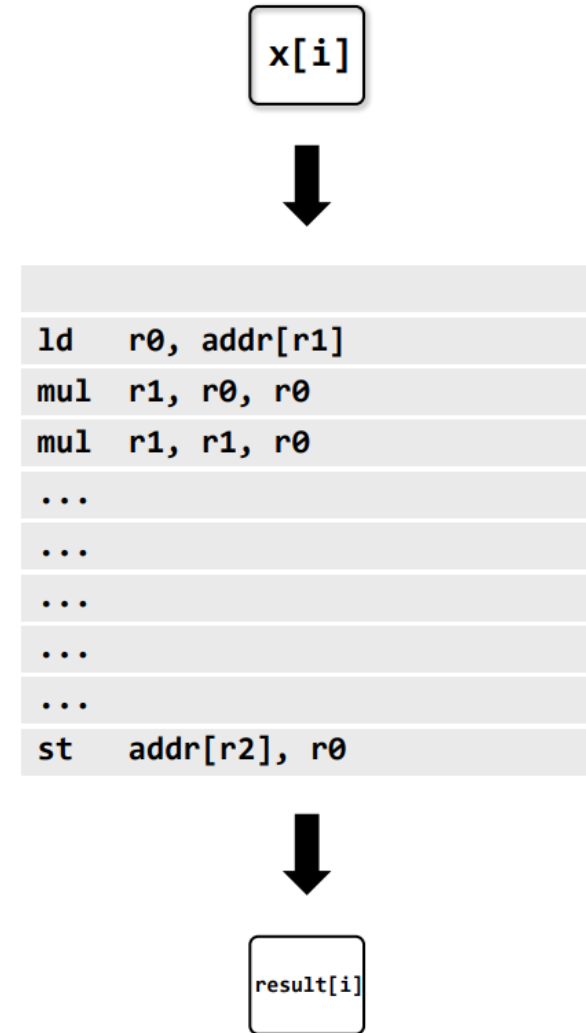



```
ld    r0, addr[r1]
mul   r1, r0, r0
mul   r1, r1, r0
...
...
...
...
...
st    addr[r2], r0
```

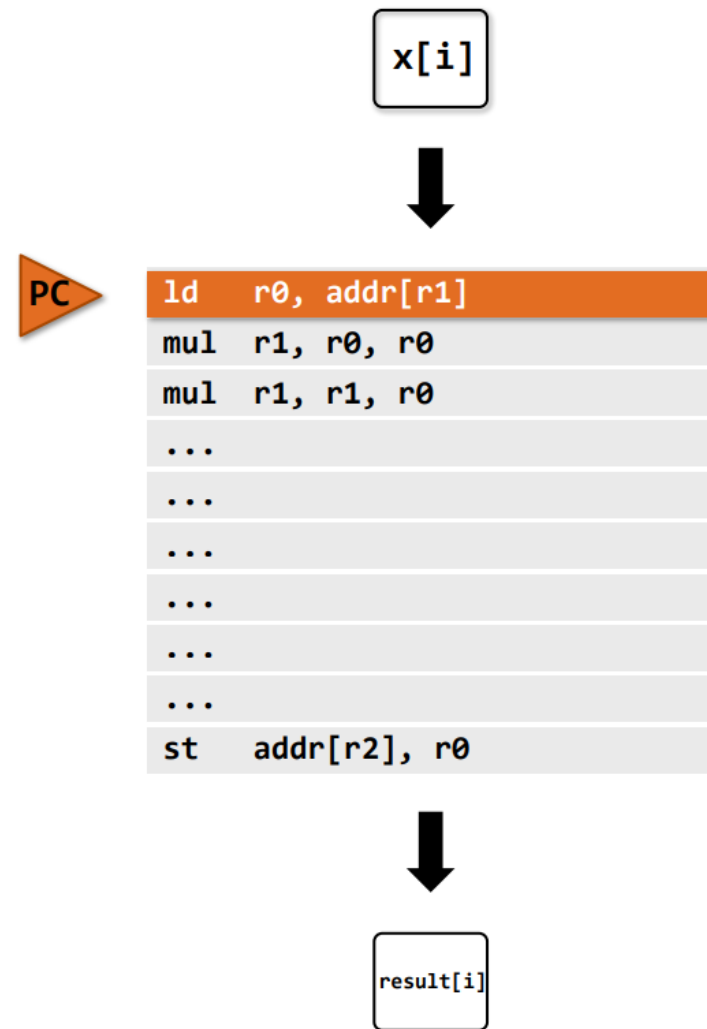
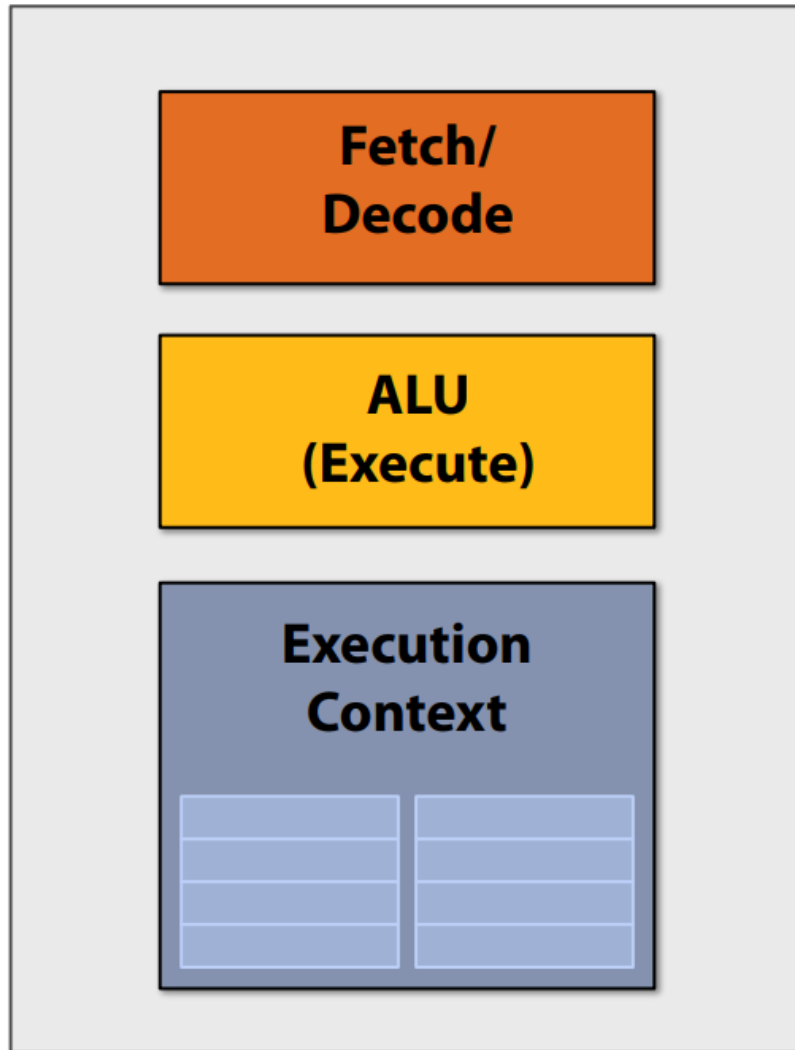


result[i]

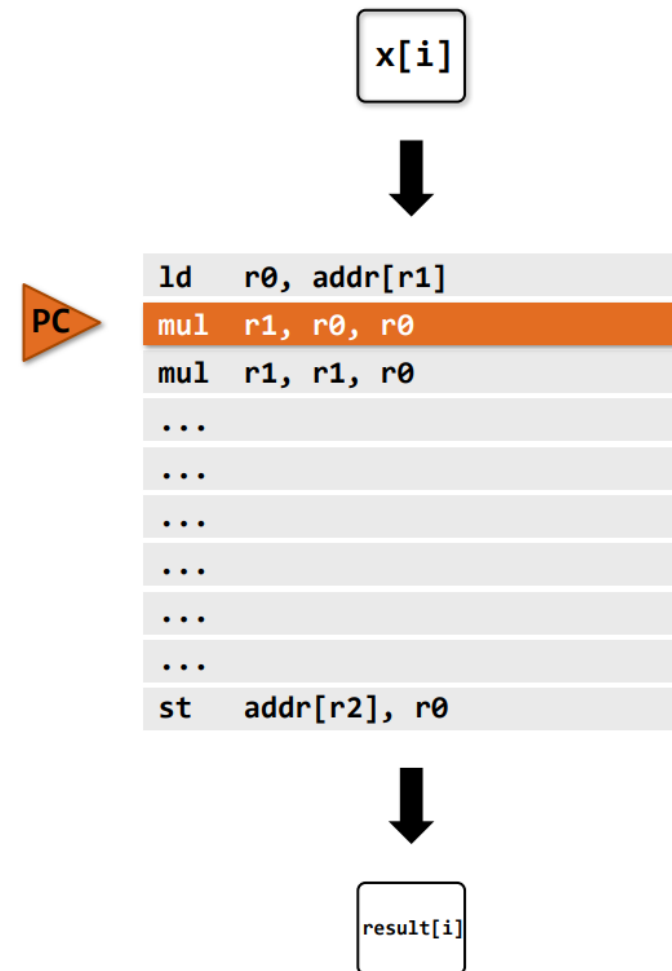

Execute Program



Execute One Instruction per clock



Execute One Instruction per clock



Von Neumann Bottleneck

Limited throughput due to data transfer between CPU and memory

Improvements to von Neumann to solve the bottleneck or to make CPUs faster

Caching

Virtual memory

Instruction-level parallelism

Thread-level parallelism

Data-level parallelism

Improvements to von Neumann

Caching

Virtual memory

Instruction-level parallelism

Thread-level parallelism

Data-level parallelism

Limitations of Memory System Performance

Memory system performance is largely captured by two parameters, latency and bandwidth

Latency: the time from the issue of a memory request to the time the data is available at the processor

Bandwidth: the rate at which data can be pumped to the processor by the memory system

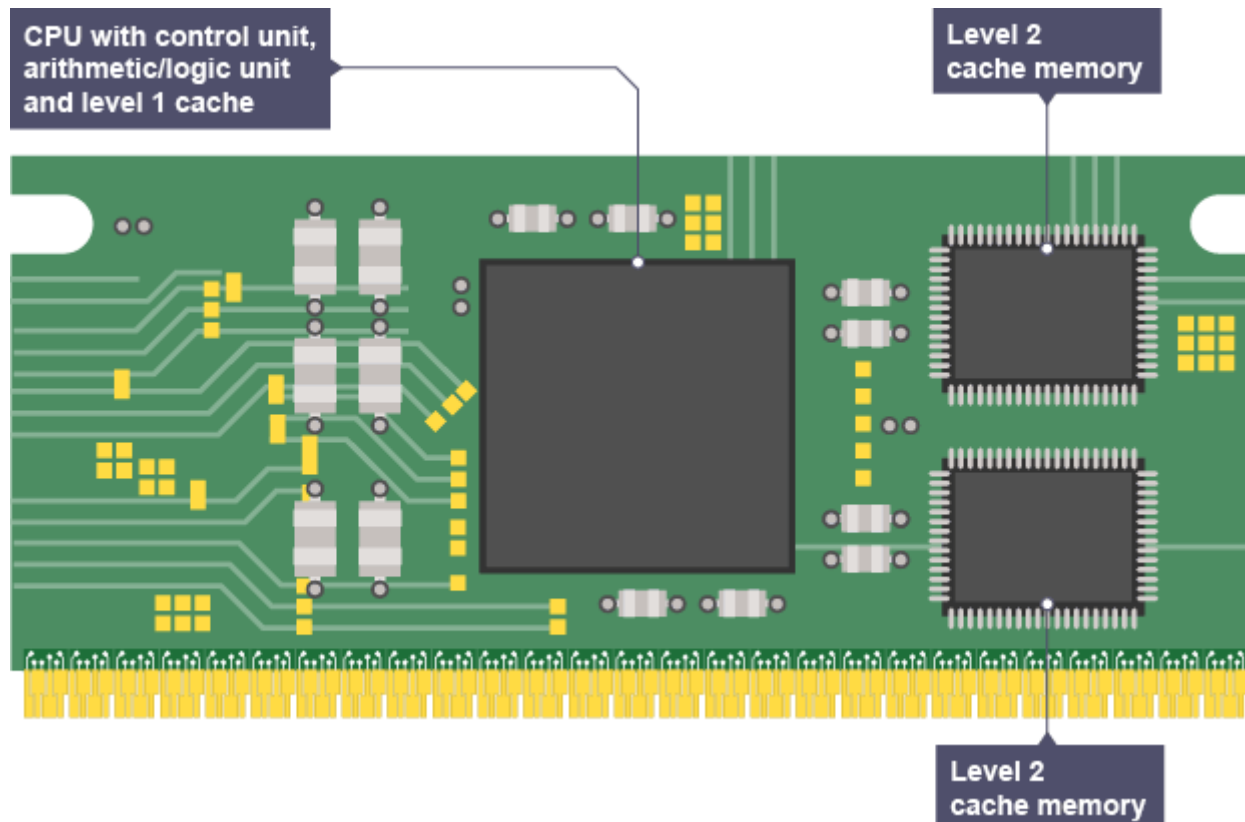
Latency vs bandwidth: Fire-hose analogy

If the water comes out of the hose 2 seconds after the hydrant is turned on, the **latency** of the system is 2 seconds (put out a fire immediately)

If the hydrant delivers water at the rate of 1 gallon/second, the **bandwidth** of the system is 1 gallon/second (large fires)

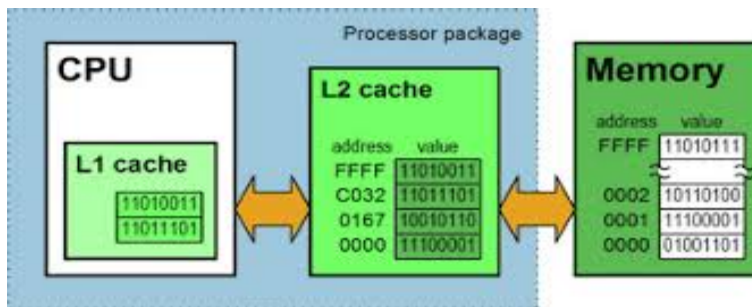
Caching

Caches: small and fast memory elements between the processor and main memory



Caching

Collection of memory locations acting as a low-latency high-bandwidth storage



CPU can access data in caches more quickly than it can access main memory

Which data stored in caches?

Locality



Temporal locality

Recently accessed items will be accessed in the near future (e.g., code in loops)

Spatial locality

Items at addresses close to the addresses of recently accessed items will be accessed in the near future (e.g., elements of an array)

```
float z[1000];  
...  
sum = 0.0;  
for (i = 0; i < 1000; i++)  
    sum += z[i];
```



Improvements to von Neumann

Caching

Virtual memory

Instruction-level parallelism

Thread-level parallelism

Data-level parallelism

Virtual Memory

If we run a program that accesses very large datasets, all data may not fit into main memory

Virtual memory

Main memory as a cache for secondary storage

Keeping only the active parts of a program in main memory, keeping idle parts in a block of secondary storage called **swap space**

Blocks of data: **pages**, virtual page numbers for programs

Mapping virtual page numbers to physical addresses: **page table**

Improvements to von Neumann

Caching

Virtual memory

Instruction-level parallelism

Thread-level parallelism

Data-level parallelism

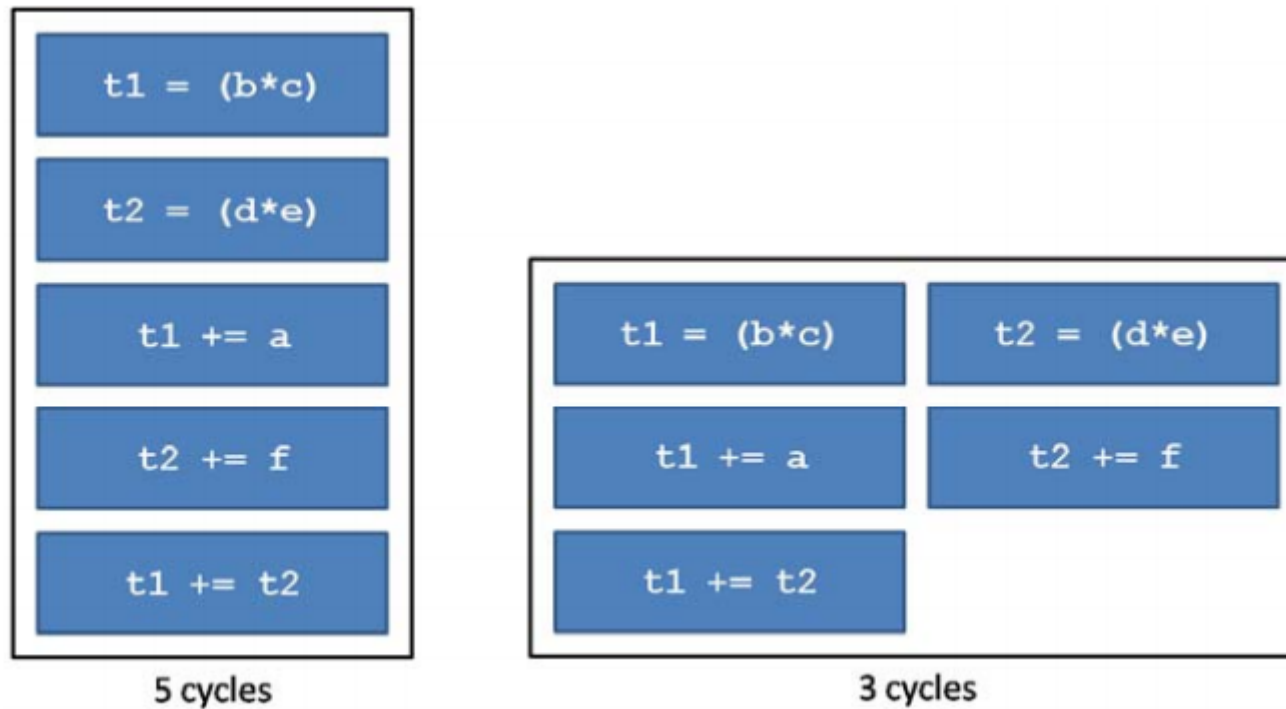
Instruction-Level Parallelism (ILP)

Having multiple processor components or functional units simultaneously executing instructions

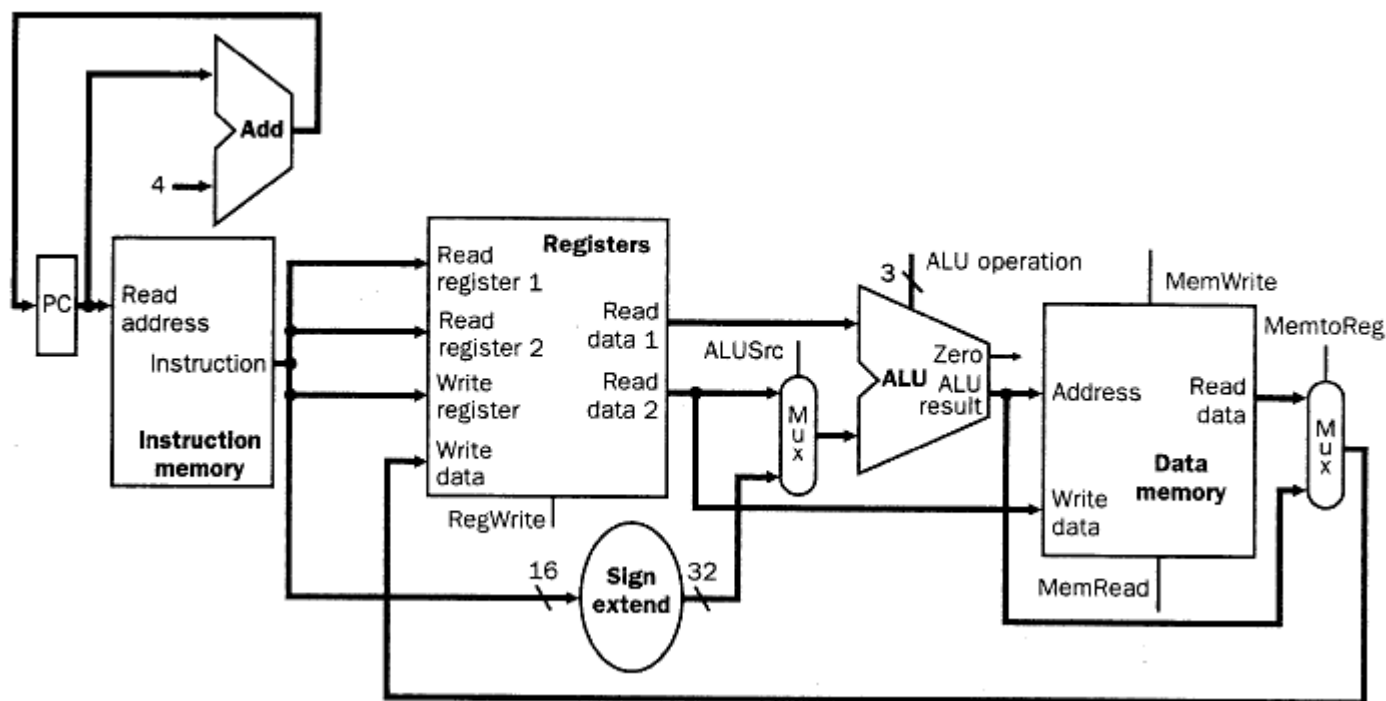
Pipelining

Multiple issue

Instruction-Level Parallelism





Single-cycle CPU Datapath

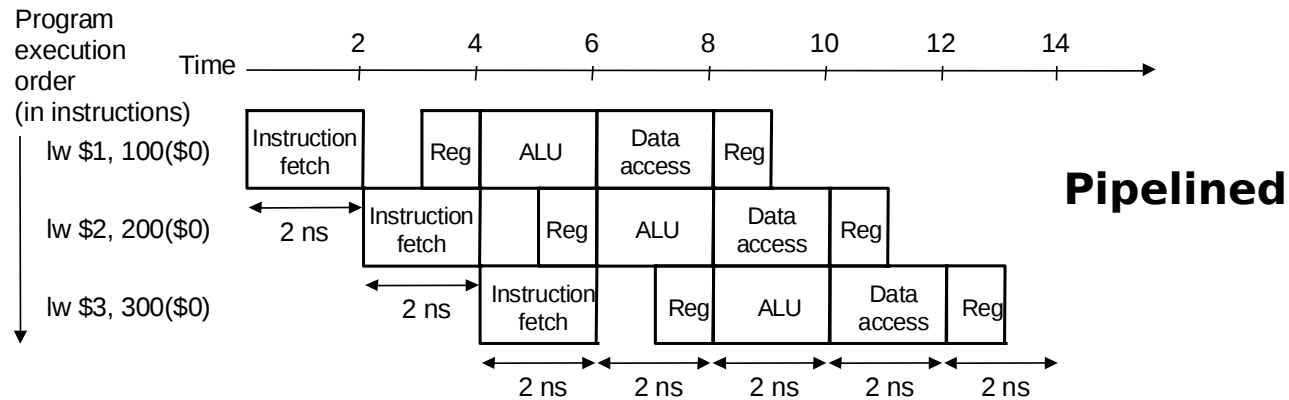
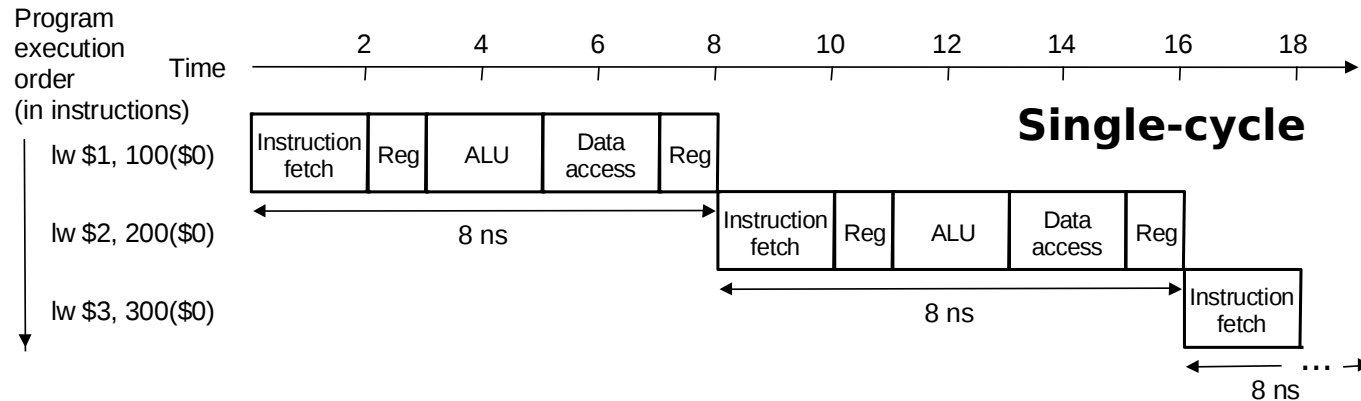


Pipelining

Takes individual pieces of hardware or functional units and connects them in sequence



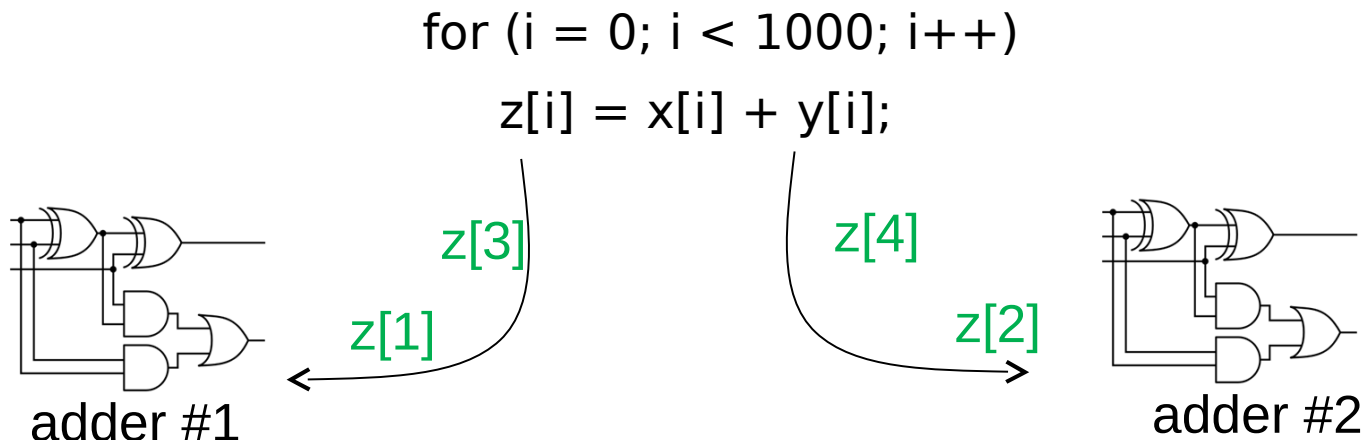
Pipelined vs Single-Cycle Execution



Multiple Issue Processors

Replicate functional units and try to simultaneously execute multiple different instructions in a clock cycle

If we have two adders, we can halve the time it takes to execute the loop (first computes $z[0]$, the second computes $z[1]$, and so on)



Multiple Issue Processors

Superscalar

Static

Dynamic

Speculative

VLIW

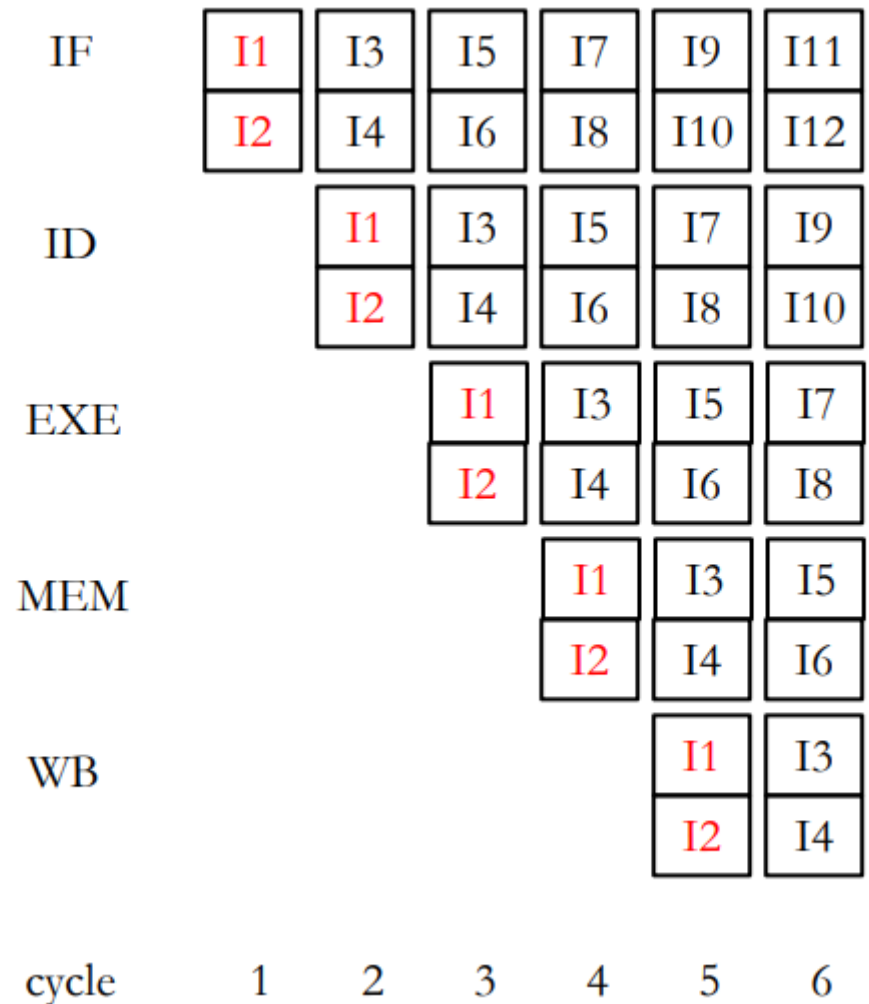
Superscalar Processors

Start two instructions per clock cycle by scheduling functional units

Static: If functional units are scheduled at compile-time

Dynamic: If they are scheduled at run-time

Speculative: Scheduling by making some predictions about the outcome of the instructions




Speculation

In order to make use of multiple issue, the system must find instructions that can be executed simultaneously

In speculation, the compiler or the processor makes a guess about an instruction, and then executes the instruction on the basis of the guess

Speculation

```
z = x + y ;  
if ( z > 0 )  
    w = x + 2 ;  
else  
    w = y * 3;
```



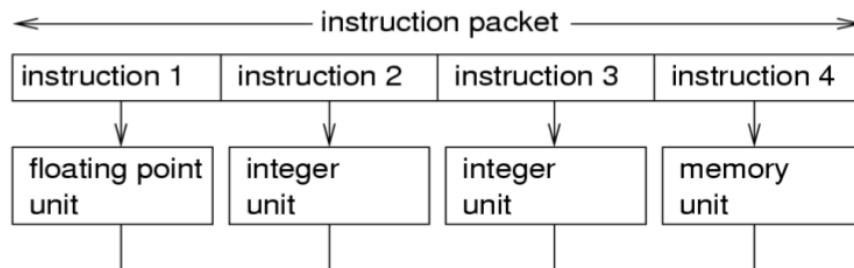
If the system speculates incorrectly,
it must go back and recalculate $w = y * 3$

VLIW (Very large instruction window)

Uses multiple, independent functional units

Does not issue multiple independent instructions to the units

Packages the multiple operations into one very long instruction



Improvements to von Neumann

Caching

Virtual memory

Instruction-level parallelism

Thread-level parallelism

Data-level parallelism

Thread-Level Parallelism

Parallelism through the simultaneous execution of different threads

Multiple data from multiple instructions (MIMD)

Multiprocessors

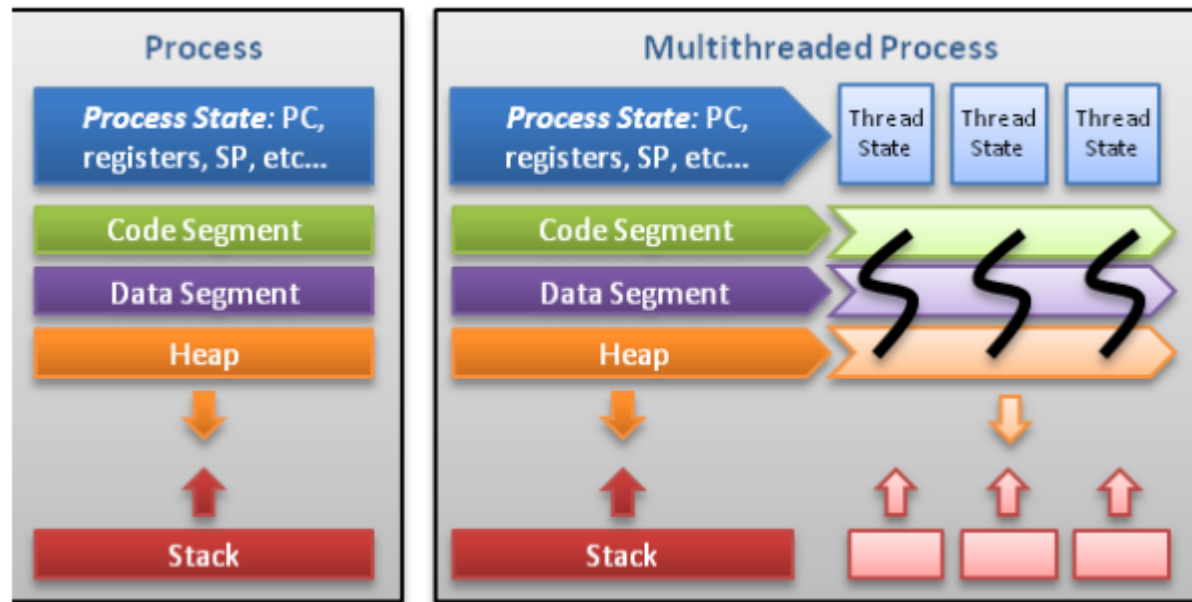
Shared-memory multiprocessors

Distributed-memory multiprocessors

Threads

Threads use, and exist within, the process resources

Scheduled and run as independent entities




Thread as Function Instance

```
for (i = 0; i < n; i++)  
    y[i] = dot_product(row(A, i), b);
```




```
for (i = 0; i < n; i++)  
    y[i] = create_thread(dot_product(row(A, i), b));
```



Shared-Memory System

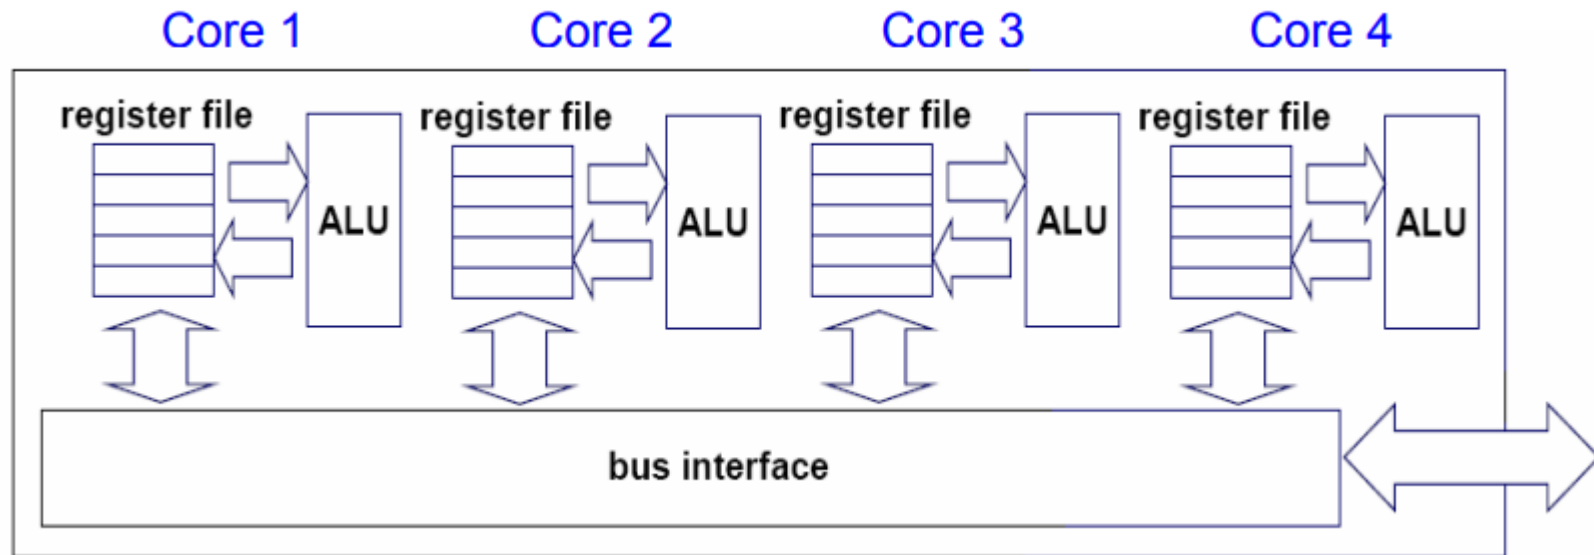
A collection of processors connected to a memory system via an interconnection network, each processor can access each memory location



Multicore Processors

Multiple CPUs on a single chip

The most widely available shared-memory systems use one or more multicore processors



Shared-Memory MPs

Load/store instructions can directly access all part of the shared address space

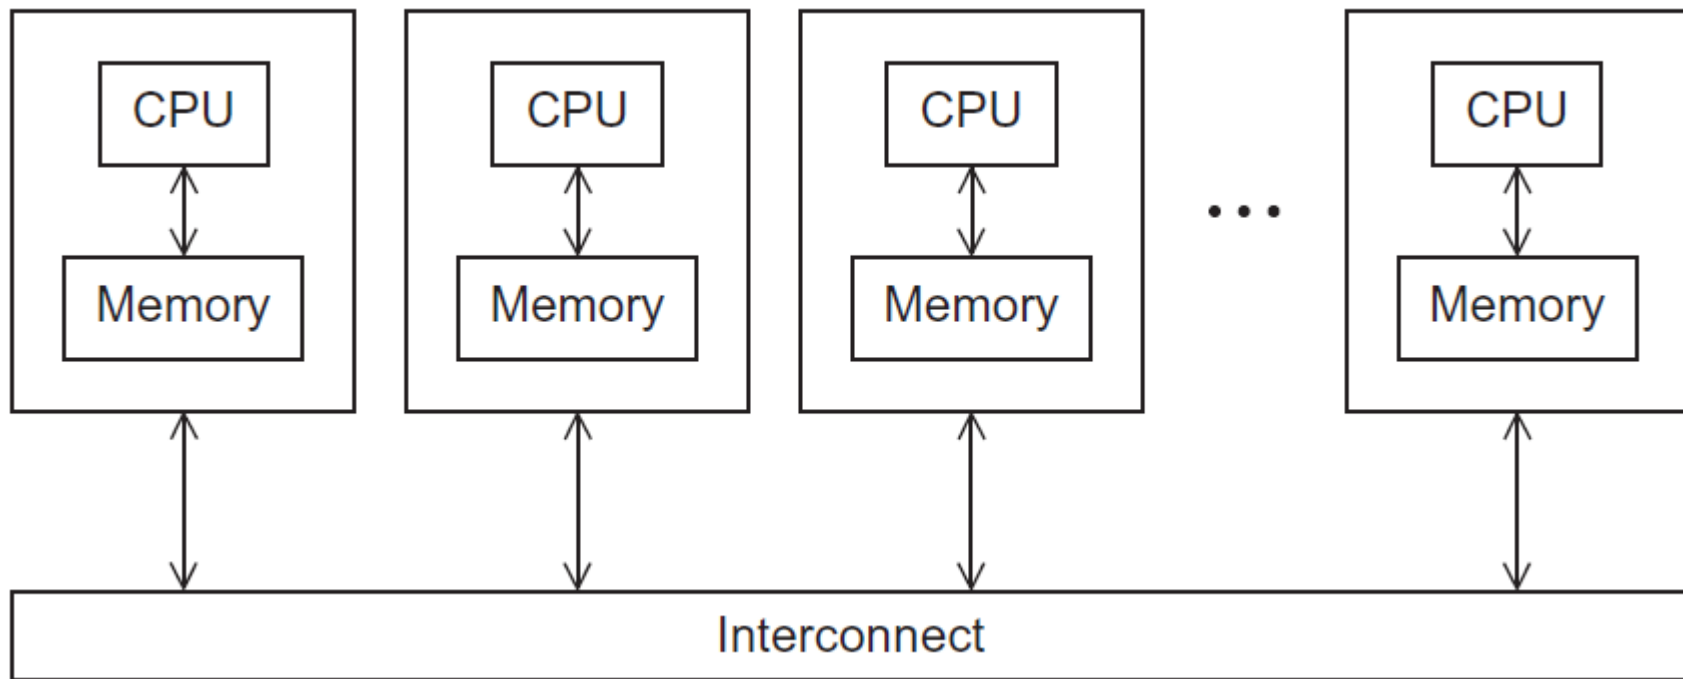
The same instruction supported by uniprocessors

Simple and fast mechanism to communicate and share information

Specific instructions to build mechanisms of synchronization

Distributed-Memory System

A collection of computers connected via an interconnection network with their own memory, processors need to communicate each other by *message passing*



Improvements to von Neumann

Caching

Virtual memory

Instruction-level parallelism

Thread-level parallelism

Data-level parallelism

Data-Level Parallelism


Single instruction processing multiple data (SIMD)

Vector processors

SIMD Multimedia instructions

GPUs

Data Parallelism - Vector Addition Example



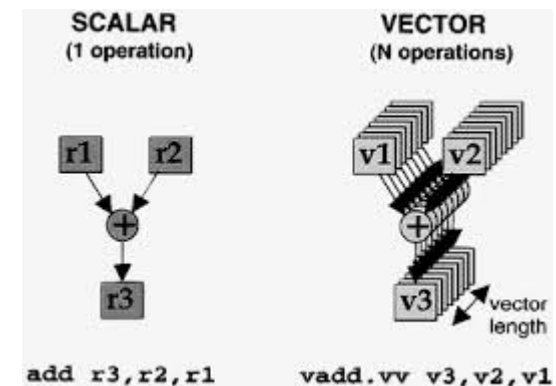
Vector Processors

Operate on arrays or vectors of data while conventional CPUs operate on individual data elements or scalars


Vector registers

Vectorized and pipelined functional units

Vector instructions



Vector Architecture



AVX-512

Advanced vector extensions by Intel 512-bit extensions

SIMD Multimedia Instructions


Since multimedia applications require narrower data types than 64-bit processors, partitioning registers and datapath

Treat a 64-bit register as a vector of two 32-bit or four 16-bit or eight 8-bit values (like short vectors)

Partition 64-bit datapaths to handle multiple narrow operations in parallel


Graphics Processing Units (GPUs)

Graphics workloads: Identical, independent, streaming computation on pixels



Graphics Processing Units (GPUs)

A massively parallel architecture consisting of thousands of smaller, more efficient cores designed for handling multiple tasks simultaneously



Graphics Pipeline

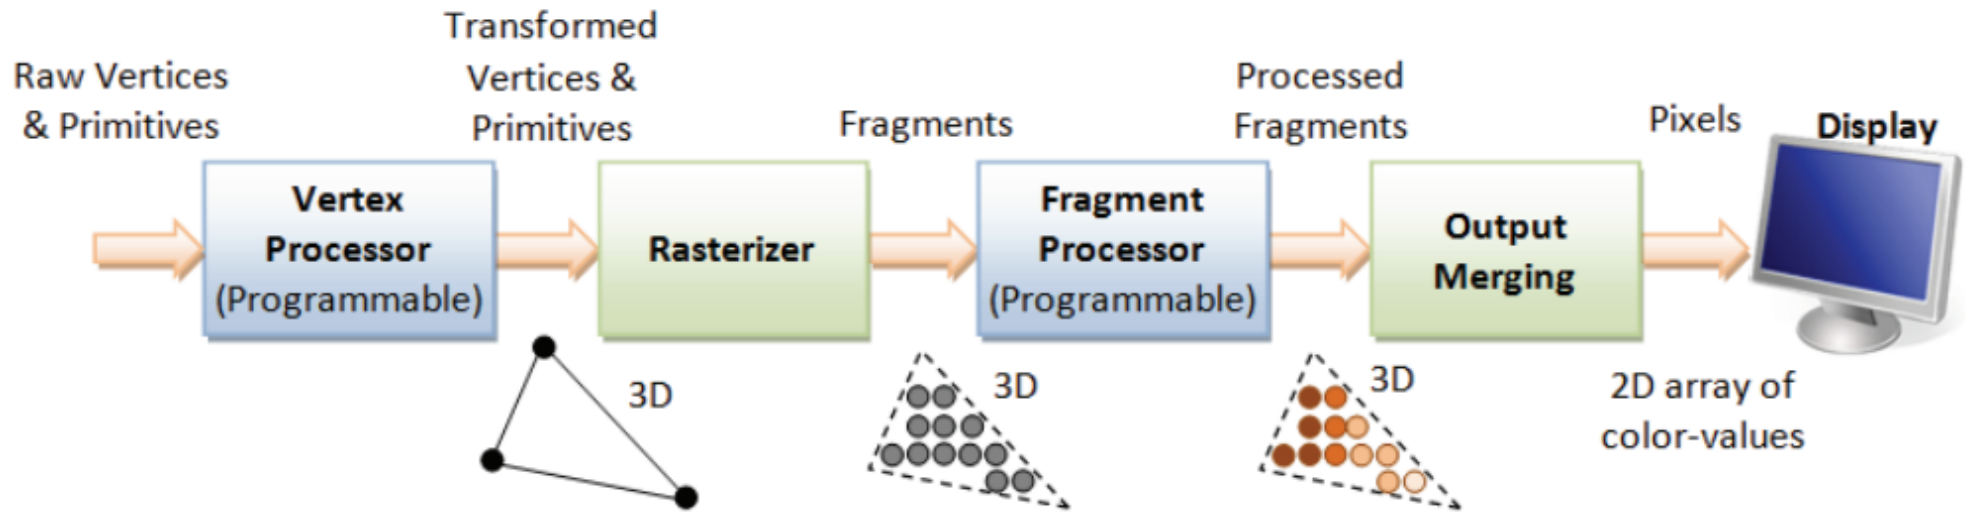
Due to the very nature of the computer graphics, most graphic pipelines have been structured as a computing states with data flowing as streams between them

Converts the internal representation into an array of pixels that can be sent to a screen

Several stages of this pipeline (shader functions) are programmable

Shader functions are processed independently, but no explicit parallel programming, they are implicitly parallel

Graphics Pipeline



1. **Vertex Processing:** Process and transform individual vertices & normals.
2. **Rasterization:** Convert each primitive (connected vertices) into a set of fragments. A fragment can be treated as a pixel in 3D spaces, which is aligned with the pixel grid, with attributes such as position, color, normal and texture.
3. **Fragment Processing:** Process individual fragments.
4. **Output Merging:** Combine the fragments of all primitives (in 3D space) into 2D color-pixel for the display.

Evolution of Graphics Pipelines

Fixed-function pipelines

configurable but not programmable

DirectX, OpenGL, early NVIDIA GeForce GPUs

First step toward achieving true general shader programmability

GeForce 3, 2001

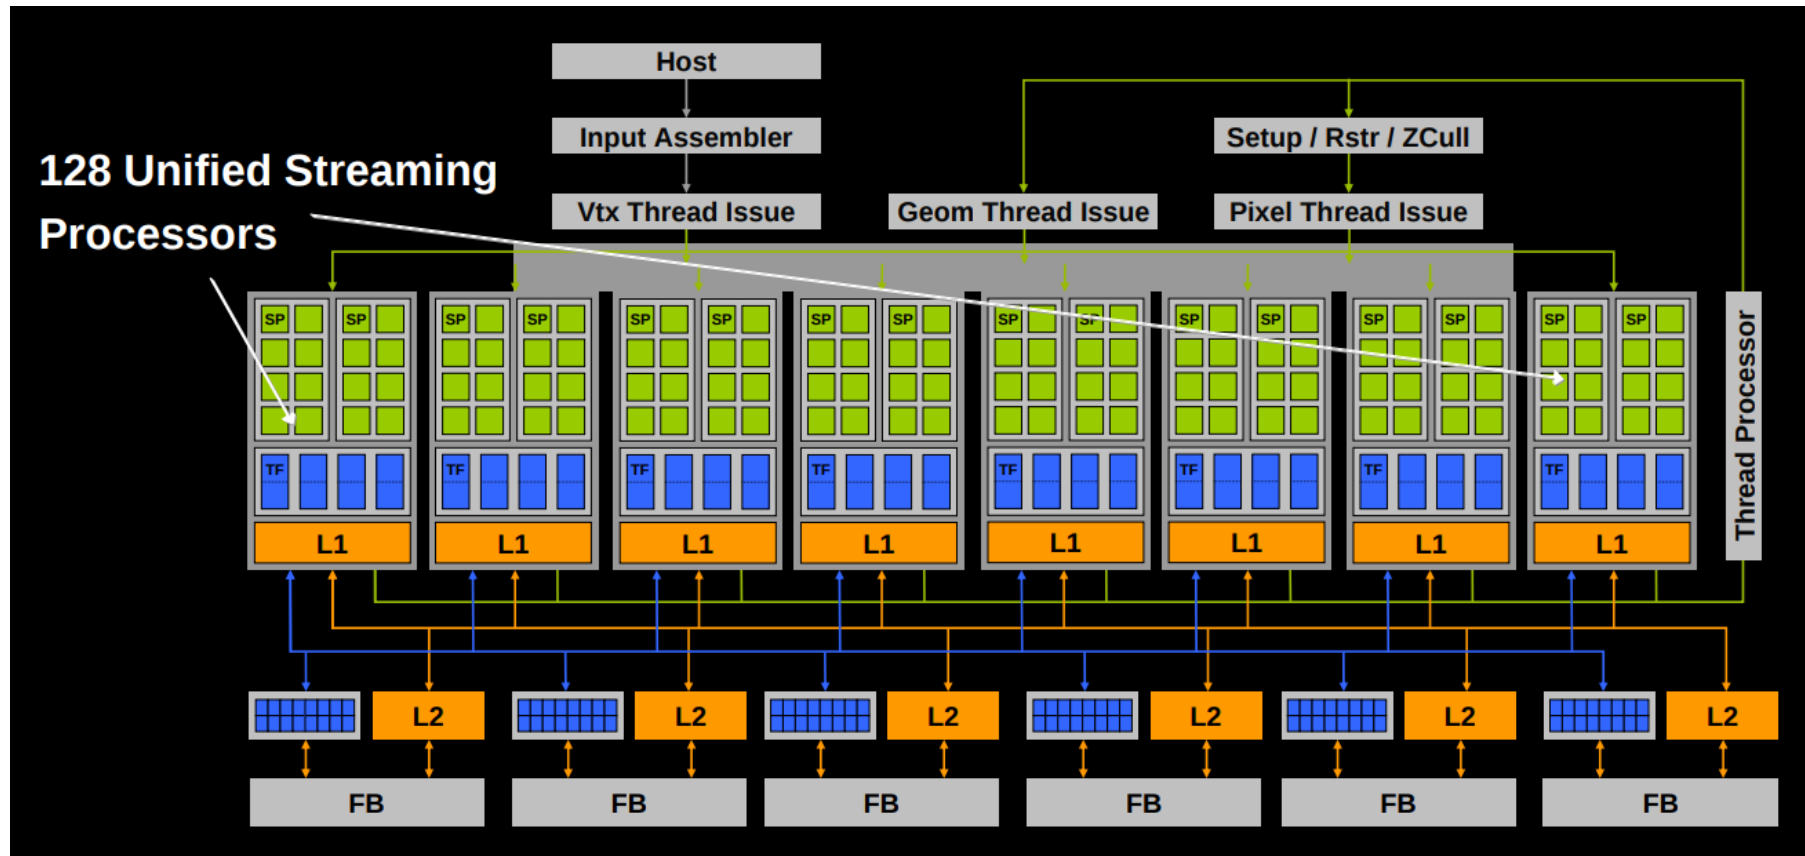
Mapping the separate **programmable graphics stages to an array of unified processors**

GeForce 8800 GPU (G80), 2006

GPGPU

GPU like a preprocessor, resembled high-performance parallel computers

G80 Replaces Pipeline Model



Inside a Stream Multiprocessor (SM)

Scalar register-based ISA

Multithreaded Instruction Unit

Up to 1024 concurrent threads

Hardware thread scheduling

In-order issue

8 SP: Thread Processors

IEEE 754 32-bit floating point

32-bit and 64-bit integer

16K 32-bit registers

2 SFU: Special Function Units


sin, cos, log, exp

Double Precision Unit

IEEE 754 64-bit floating point

Fused multiply-add

16KB Shared Memory



GPU Computing Today

GPUs and CUDA bring parallel computing to the masses

Over 100M CUDA-capable GPUs sold to date

60K CUDA developers

A “developer kit” (i.e. GPU) costs ~\$200 (for 500 GFLOPS)

Data-parallel supercomputers are everywhere!

CUDA makes this power accessible

We're already seeing innovations in data-parallel computing

Massively parallel computing has become a commodity technology!