

**CENG443**

# **Heterogeneous Parallel Programming**

## **CUDA Performance**

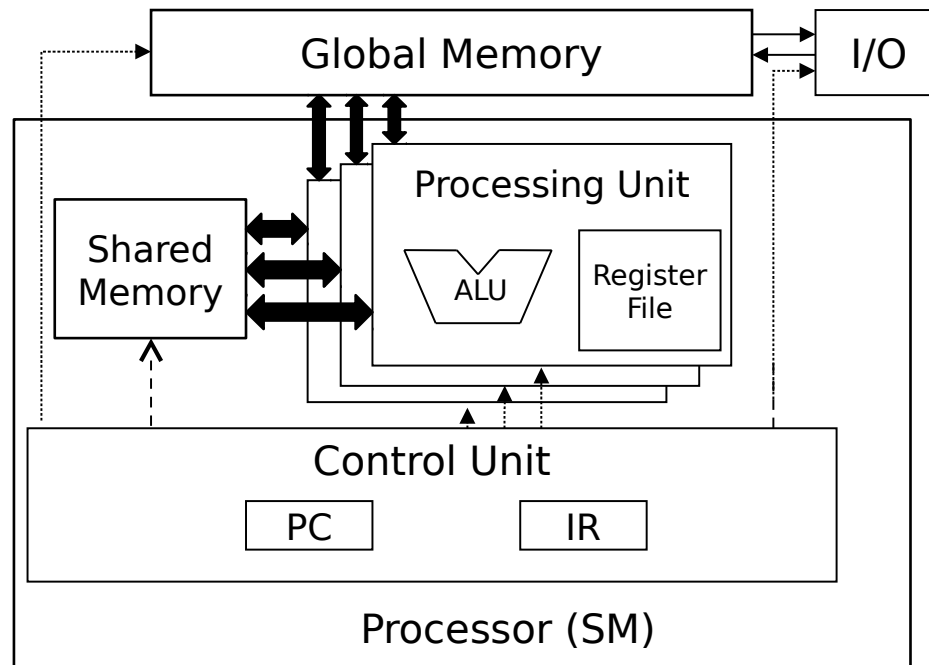


# CUDA Kernel Performance

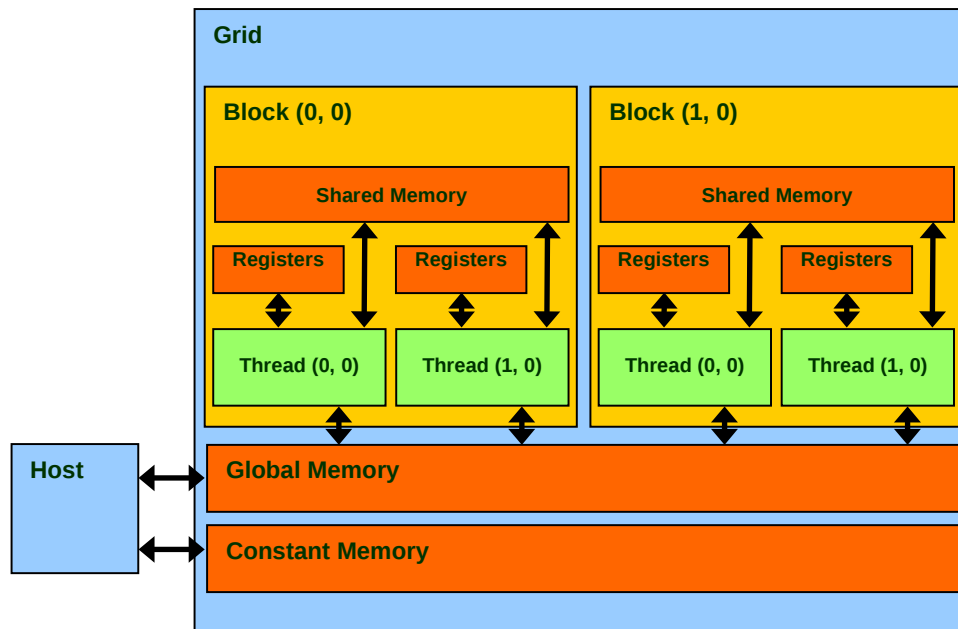
**Global memory bandwidth**

**Warp scheduling**

# Hardware View of CUDA Memories



# Global Memory



- Typically implemented in DRAM
- High access latency:  
400-800 cycles
- Finite access bandwidth
- Potential of traffic congestion
- Throughput up to 177GB/s

# Global Memory

**GPU device with the global memory bandwidth  
1,000 GB/s, or 1 TB/s**

**4 bytes in each single-precision floating-point value**

**$1000/4=250$  giga single-precision operands per  
second to be loaded**

no more than 250 giga floating-point operations per second  
(GFLOPS)

**Peak single-precision performance of 12 TFLOPS**

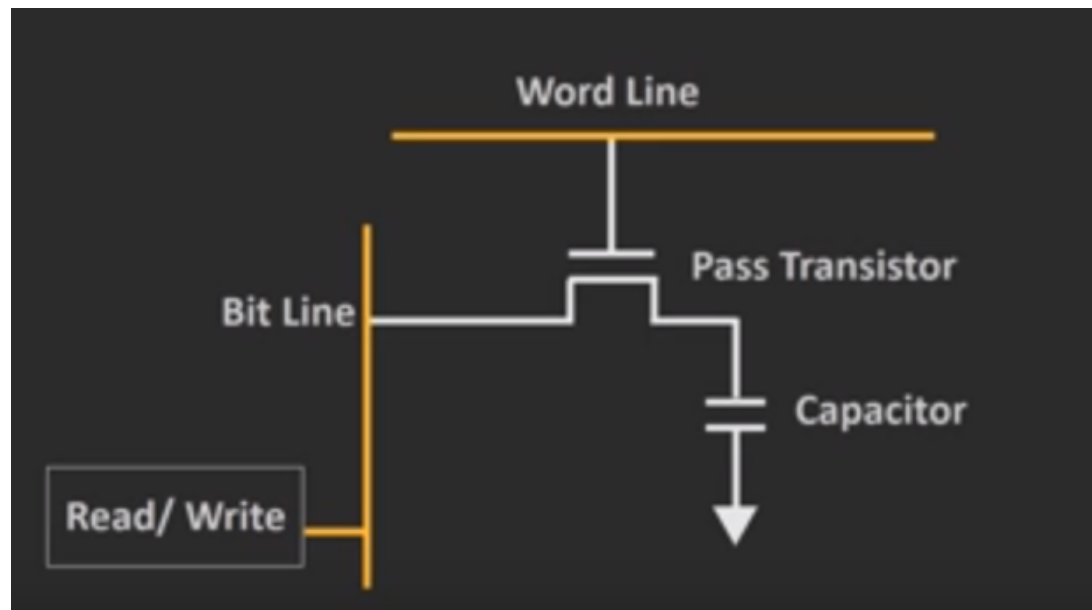
**Tiny fraction =  $250/12000 = 2\%$**

**Need to find ways of reducing global memory  
accesses!**

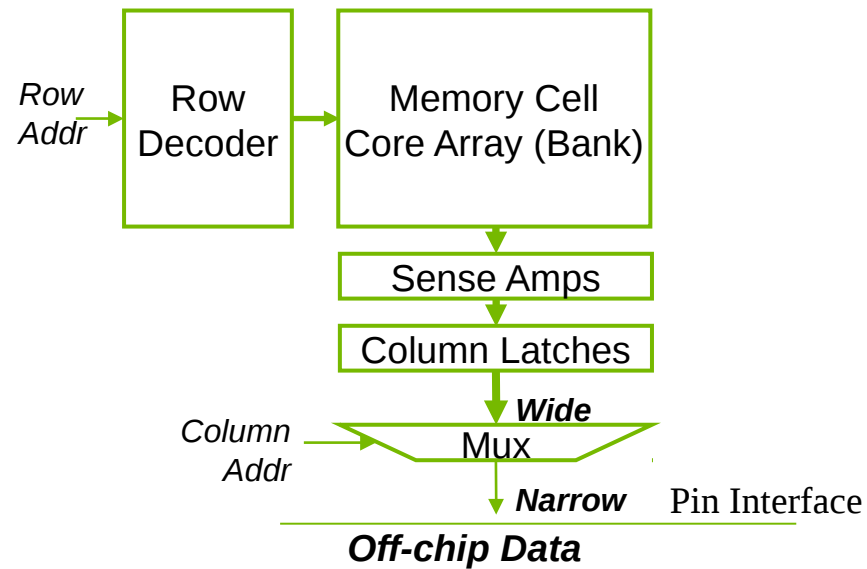
# DRAM (Dynamic Random Access)

**Single bit is stored as a charge in a capacitor, one transistor is used for access**


Whether the capacitor is charged or discharged indicates storage of 1 or 0



# DRAM Bank Organization



# DRAM Bank Operation





# DRAM Burst


**Each time a DRAM location is accessed, a range of consecutive locations that includes the requested location are actually accessed**

**The data from all these consecutive locations can be transferred at very high-speed to the processor**

**These consecutive locations accessed and delivered are referred to as DRAM *bursts***

**If an application makes focused use of data from these bursts, the DRAMs can supply the data at a much higher rate than if a truly random sequence of locations were accessed**

# DRAM Burst-A System View



## Each address space is partitioned into burst sections

Whenever a location is accessed, all other locations in the same section are also delivered to the processor

## Basic example: a 16-byte address space, 4-byte burst sections

In practice, we have at least 4GB address space, burst section sizes of 128-bytes or more

# Memory Coalescing

**Recognizing the burst organization of modern DRAMs, current CUDA devices employ a technique that allows the programmers to achieve high global memory access efficiency by organizing memory accesses of threads into favorable patterns**

**Takes advantage of the fact that threads in a warp execute the same instruction at any given point in time**

# Memory Coalescing


**When all threads in a warp execute a load instruction, the hardware detects whether they access consecutive global memory locations**

**The most favorable access pattern is achieved when all threads in a warp access consecutive global memory locations**

**The hardware combines, or *coalesces*, all these accesses into a consolidated access to consecutive DRAM locations**


**Such coalesced access allows the DRAMs to deliver data as a burst**

# Memory Coalescing




**When all threads of a warp execute a load instruction, if all accessed locations fall into the same burst section, only one DRAM request will be made and the access is fully coalesced**

# Coalesced Accesses




All threads participate



Some Threads Do Not Participate

# Uncoalesced Accesses



**When the accessed locations spread across burst section boundaries:**


- Coalescing fails

- Multiple DRAM requests are made


- The access is not fully coalesced

**Some of the bytes accessed and transferred are not used by the threads**

# Coalesced Access



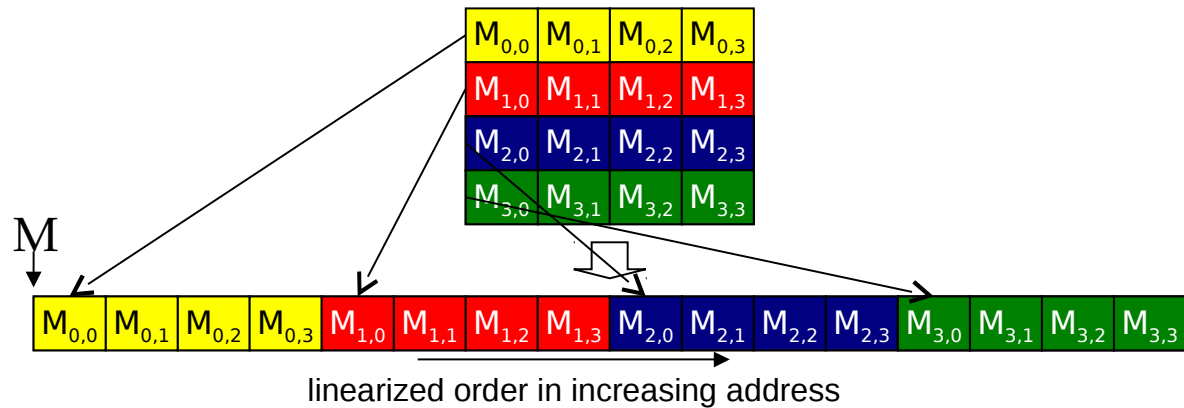
All threads participate




Some Threads Do Not Participate



# A 2D C Array in Linear Memory Space




# Matrix Multiplication



# A Basic Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, int Width) {  
    // Calculate the row index of the P element and M  
    int Row = blockIdx.y*blockDim.y+threadIdx.y;  
  
    // Calculate the column index of P and N  
    int Col = blockIdx.x*blockDim.x+threadIdx.x;  
  
    if ((Row < Width) && (Col < Width)) {  
        float Pvalue = 0;  
        // each thread computes one element of the block sub-matrix  
        for (int i = 0; i < Width; ++i) {  
            Pvalue += M[Row*Width+i]*N[i*Width+Col];  
        }  
        P[Row*Width+Col] = Pvalue;  
    }  
}
```

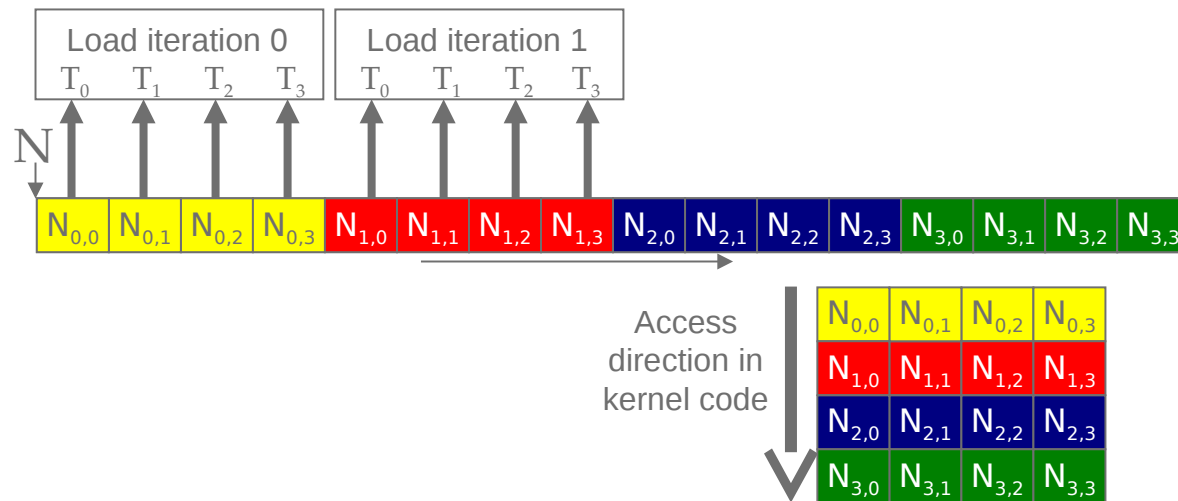
# Matrix Multiplication in C 2D arrays



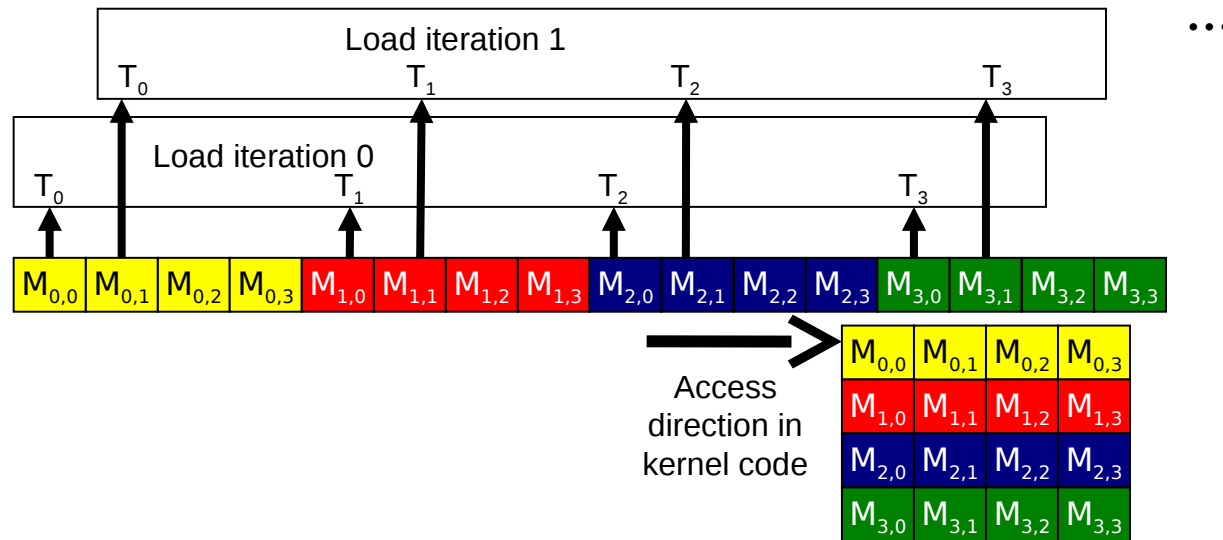
$i$  is the loop counter in the inner product loop of the kernel code

$\text{Col} = \text{blockIdx.x} \times \text{blockDim.x} + \text{threadIdx.x}$   
 $\text{Row} = \text{blockIdx.y} \times \text{blockDim.y} + \text{threadIdx.y}$

# N Accesses are Coalesced



# M Accesses are not Coalesced




# Using Shared Memory for Coalescing

If an algorithm requires a kernel code to iterate through data along the row direction, one can use the shared memory to enable memory coalescing-tiles are loaded in a coalesced pattern

Have each thread load an M element and an N element at the same relative position as its P element.

```
int tx = threadIdx.x  
int ty = threadIdx.y  
Accessing tile 0 2D indexing:  
M[Row][tx]  
N[ty][Col]
```



# Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, Int Width)
{
    __shared__ float ds_M[TILE_WIDTH][TILE_WIDTH];
    __shared__ float ds_N[TILE_WIDTH][TILE_WIDTH];

    int bx = blockIdx.x; int by = blockIdx.y;
    int tx = threadIdx.x; int ty = threadIdx.y;

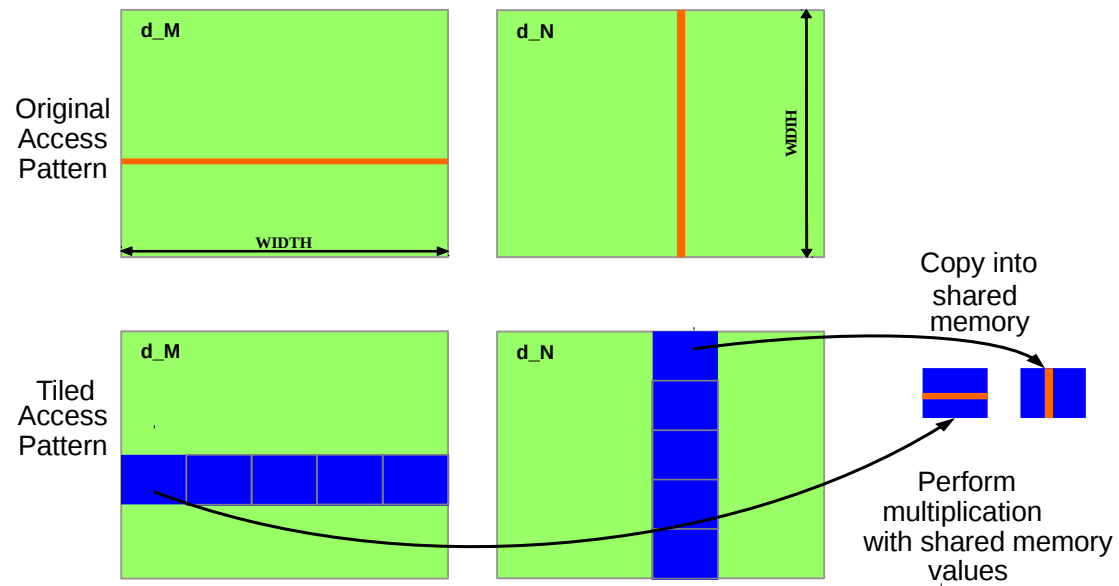
    int Row = by * blockDim.y + ty;
    int Col = bx * blockDim.x + tx;
    float Pvalue = 0;

    // Loop over the M and N tiles required to compute the P element
    for (int p = 0; p < Width/TILE_WIDTH; ++p) {
        // Collaborative loading of M and N tiles into shared memory
        ds_M[ty][tx] = M[Row*Width + p*TILE_WIDTH+tx];
        ds_N[ty][tx] = N[(p*TILE_WIDTH+ty)*Width + Col];
        __syncthreads();

        for (int i = 0; i < TILE_WIDTH; ++i) Pvalue += ds_M[ty][i] * ds_N[i][tx];
        __syncthreads();
    }
    P[Row*Width+Col] = Pvalue;
}
```

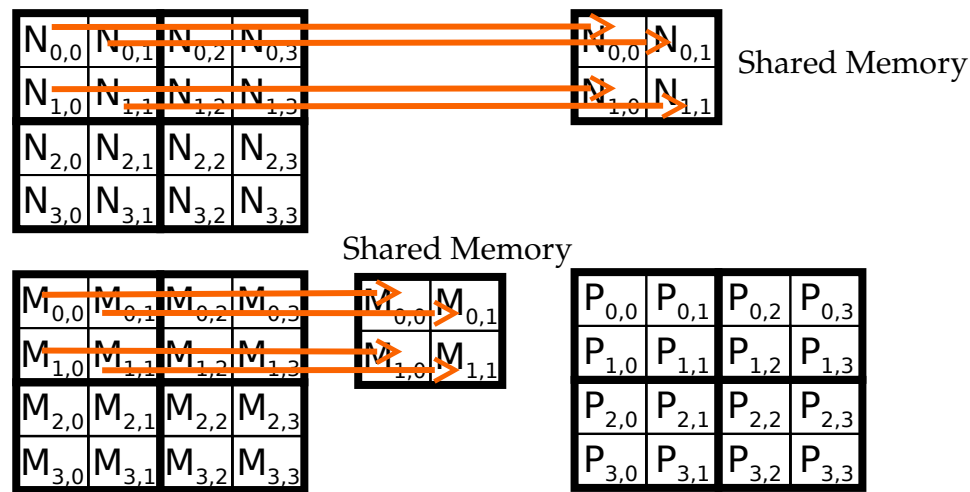


# Corner Turning



# Phase 0 Load for Block (0,0)

Each thread loads one M element and one N element in tiled code



Turned a vertical access pattern into a horizontal access pattern

# Tiled Algorithm


**Loads to both the M and N elements are coalesced**

**1) number of memory loads is reduced due to the reuse of data in the shared memory**


**2) the remaining memory loads are coalesced so the DRAM bandwidth utilization is further improved**

**Tiled kernel can run 30x faster than the simple kernel**

# Multiple DRAM Banks



# Channels and Banks in DRAM



# DRAM Bursting with Banking



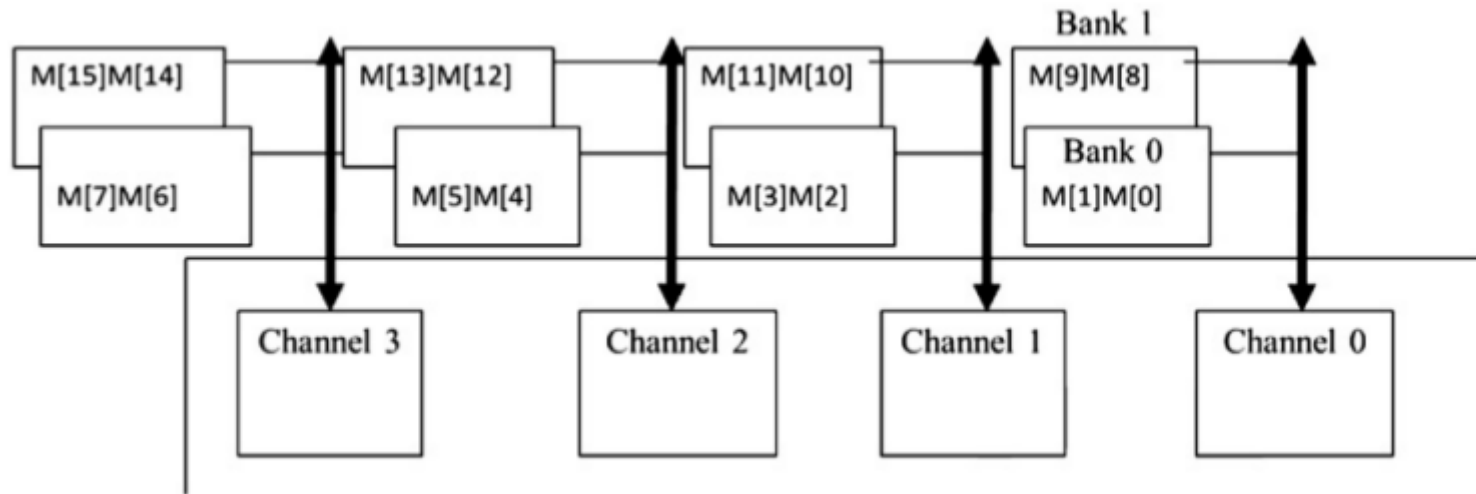
Single-Bank burst timing, dead time on interface




Multi-Bank burst timing, reduced dead time

# Array Elements into Channels and Banks

## Interleaved data distribution



# Matrix Multiplication




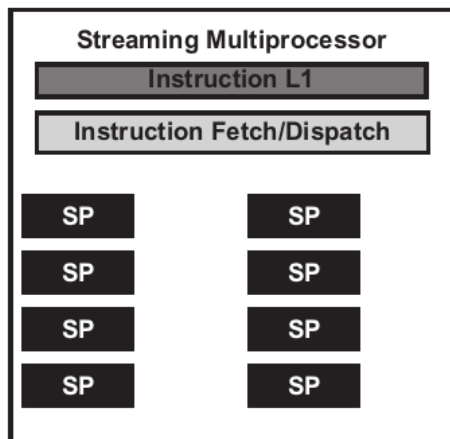



# M Elements


Tiles loaded by	Block 0,0	Block 0,1	Block 1,0	Block 1,1
Phase 0 (2D index)	M[0][0], M[0][1], M[1][0], M[1][1]	M[0][0], M[0][1], M[1][0], M[1][1]	M[2][0], M[2][1], M[3][0], M[3][1]	M[2][0], M[2][1], M[3][0], M[3][1]
Phase 0 (linearized index)	M[0], M[1], M[4], M[5]	M[0], M[1], M[4], M[5]	M[8], M[9], M[12], M[13]	M[8], M[9], M[12], M[13]
Phase 1 (2D index)	M[0][2], M[0][3], M[1][2], M[1][3]	M[0][2], M[0][3], M[1][2], M[1][3]	M[2][2], M[2][3], M[3][2], M[3][3]	M[2][2], M[2][3], M[3][2], M[3][3]
Phase 1 (linearized index)	M[2], M[3], M[6], M[7]	M[2], M[3], M[6], M[7]	M[10], M[11], M[14], M[15]	M[10], M[11], M[14], M[15]

# Warps

**Each block is divided into 32-thread warps**  
**Warps are scheduling units in SM**

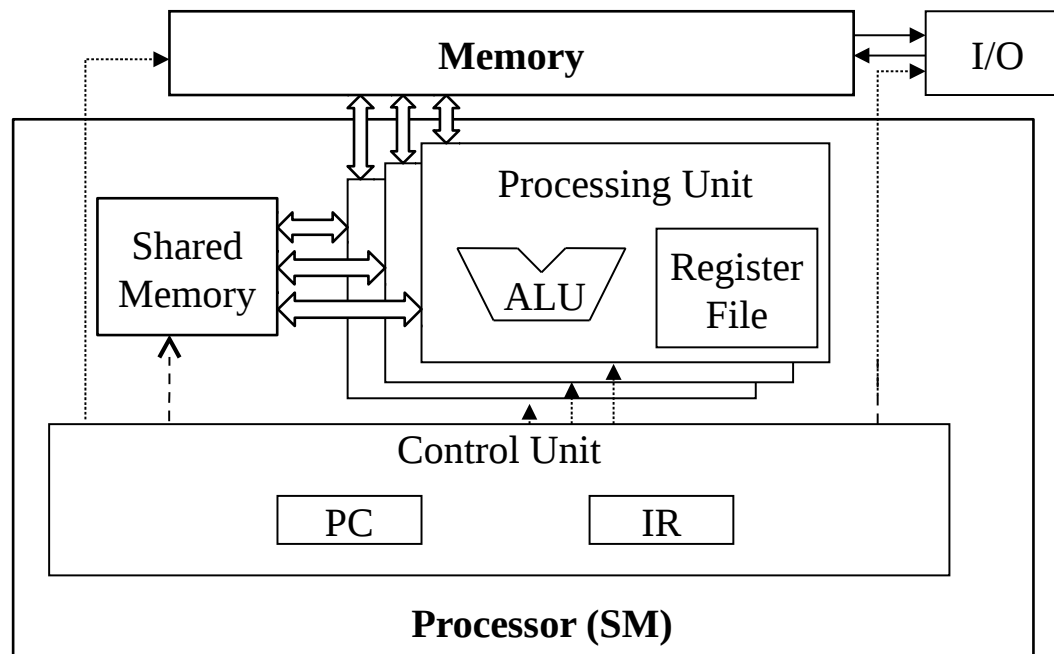


# Placing 2D Threads into Linear Order



# SMs are SIMD Processors

**Control unit for instruction fetch, decode, and control is shared among multiple processing units**



# SIMD Among Threads in a Warp

**All threads in a warp must execute the same instruction at any point in time**

**This works efficiently if all threads follow the same control flow path**

All if-then-else statements make the same decision

All loops iterate the same number of times

# Control Divergence

**Control divergence occurs when threads in a warp take different control flow paths by making different control decisions**

Some take the then-path and others take the else-path of an if-statement

Some threads take different number of loop iterations than others

**The execution of threads taking different paths are serialized in current GPUs**

The control paths taken by the threads in a warp are traversed one at a time until there is no more.

During the execution of each path, all threads taking that path will be executed in parallel

The number of different paths can be large when considering nested control flow statements

# Control Divergence Examples

**Divergence can arise when branch or loop condition is a function of thread indices**

**Example kernel statement with divergence:**

```
if (threadIdx.x > 2) { }
```

This creates two different control paths for threads in a block

Decision granularity < warp size; threads 0, 1 and 2 follow different path than the rest of the threads in the first warp

**Example without divergence:**

```
If (blockIdx.x > 2) { }
```

Decision granularity is a multiple of blocks size; all threads in any given warp follow the same path

# Example: Vector Addition Kernel

```
// Compute vector sum  $C = A + B$ 
```

```
// Each thread performs one pair-wise addition
```

```
__global__
```

```
void vecAddKernel(float* A, float* B, float* C, int n)
```

```
{
```

```
    int i = threadIdx.x + blockDim.x * blockIdx.x;
```

```
    if(i < n) C[i] = A[i] + B[i];
```

```
}
```



# Analysis for 1,000 Elements

## **Assume that block size is 256 threads**

8 warps in each block

## **All threads in Blocks 0, 1, and 2 are within valid range**

i values from 0 to 767

There are 24 warps in these three blocks, none will have control divergence

## **Most warps in Block 3 will not control divergence**

Threads in the warps 0-6 are all within valid range, thus no control divergence

## **One warp in Block 3 will have control divergence**

Threads with i values 992-999 will all be within valid range

Threads with i values of 1000-1023 will be outside valid range

## **Effect of serialization on control divergence will be small**

1 out of 32 warps has control divergence

The impact on performance will likely be less than 3%

# Example: Matrix Multiplication

## Boundary condition checks are vital for complete functionality and robustness of parallel code

The tiled matrix multiplication kernel has many boundary condition checks

The concern is that these checks may cause significant performance degradation


```
if (Row < Width && t * TILE_WIDTH + tx < Width) {  
    ds_M[ty][tx] = M[Row * Width + p * TILE_WIDTH + tx];  
} else {  
    ds_M[ty][tx] = 0.0;  
}
```

```
if (p * TILE_WIDTH + ty < Width && Col < Width) {  
    ds_N[ty][tx] = N[(p * TILE_WIDTH + ty) * Width + Col];  
} else {  
    ds_N[ty][tx] = 0.0;  
}
```

# Example: Matrix Multiplication

**Type 1. Blocks whose tiles are all within valid range until the last phase**

**Type 2. Blocks whose tiles are partially outside the valid range all the way**



**Assume 16x16 tiles and thread blocks**

**Each thread block has 8 warps (256/32)**

**Assume square matrices of 100x100**

**Each thread will go through 7 phases (ceiling of 100/16)**

**There are 49 thread blocks (7 in each dimension)**


# Type 1

**There are 42 ( $6 \times 7$ ) Type 1 blocks, with a total of 336 ( $8 \times 42$ ) warps**

**They all have 7 phases, so there are 2,352 ( $336 \times 7$ ) warp-phases**

**The warps have control divergence only in their last phase**

**336 warp-phases have control divergence**



# Type 2


**There are 7 Type 2 block assigned to load the bottom tiles, with a total of 56 ( $8*7$ ) warps**

**They all have 7 phases, so there are 392 ( $56*7$ ) warp-phases**

**The first 2 warps in each Type 2 block will stay within the valid range until the last phase**

**The 6 remaining warps stay outside the valid range**

**So, only 14 ( $2*7$ ) warp-phases have control divergence**



# Performance Impact of Control Divergence

## Type 1 Blocks:


336 out of 2,352 warp-phases have control divergence

## Type 2 Blocks:

14 out of 392 warp-phases have control divergence

**The performance impact  
is expected to be less than  
12%**

(350/2,944)



# Additional Comments

**The estimated performance impact is data dependent.**

For larger matrices, the impact will be significantly smaller

**In general, the impact of control divergence for boundary condition checking for large input data sets should be insignificant**

One should not hesitate to use boundary checks to ensure full functionality

**The fact that a kernel is full of control flow constructs does not mean that there will be heavy occurrence of control divergence**

**More control divergence for some algorithm patterns (such as parallel reduction)**